On The Impact of Atoms of Confusion in JavaScript Code

Abstract

(**castor note**) This is still too big.

Evolving software is a challenging task, particularly when the code has been poorly written or uses confusing idioms and language constructs, which might increase maintenance effort and impose a significant cognitive load on developers. For this reason, researchers have investigated possible sources of confusion in codebases, including the impact of small code patterns (hereafter atoms of confusion) that contribute to misunderstanding the source code written in languages such as C, C++, and Java. However, different languages have different features, developer communities, and development cultures. This justifies the exploration of other languages to verify whether they also exhibit confusion-inducing patterns. In this work, we fill this research gap by presenting the results of a mixed-methods research comprising two experiments, a set of interviews with practitioners, and a mining software repositories (MSR) effort. Our goal is to investigate the impact of atoms of confusion on understanding JavaScript code—a dynamically typed language whose popularity is growing in diverse application domains, e.g., web, backend, mobile, and IOT development. Both our experiments suggest that two code patterns that have been previously observed to confuse C programmers also confuse JavaScript programmers (with statistical significance, after correction): the comma operator and assignments being used as values. In addition, pre- and post-increment operators, when used as values, omitted curly braces, changes of literal encoding, and type conversions have led to confusion in at least one of the two experiments. For all these cases effect sizes were either medium or high. In addition, we have discovered that object destructuring and automatic semicolon insertion are potential JavaScript-specific atoms that have not been discussed in previous work. The interviews we conducted reinforce some of these findings while

pointing to other constructs and idioms that merit investigation in the future. In addition, our MSR effort reveals that atom candidates are frequent and used intensively in 72 popular open-source JavaScript systems. Four atom candidates appear in 90% of the analyzed projects, and two of them occur more than once for every 100 lines of code in the dataset. Our findings provide significant evidence that some code patterns should be avoided if code is expected to be understood by a wide audience. Furthermore, they motivate the creation of program transformation tools and the extension of linters to address these potentially harmful idioms in JavaScript codebases.

1. Introduction

Developers are often confused while reading unfamiliar code [1]. According to Hermans [2], that confusion can stem from three sources: (i) lack of knowledge, i.e., not knowing what an element in the program does; (ii) lack of information, i.e., not knowing how a program element works; and (iii) lack of processing power, i.e., the inability to combine all the elements in a program and make sense of its execution in one's head. Even small pieces of code can be confusing for developers [3, 4] at first glance.

For instance, consider the pair of JavaScript code snippets in Listing 1 (adapted from NERVJS/TARO). The code snippet on the left-hand side is arguably harder to understand than the one on the right-hand side. The former uses the logical AND operator (`&&`) beyond its lexical meaning [5] to determine the control flow of the program. Furthermore, the expression has a potential side effect, due to the use of the post-increment operator (`s++`). The snippet on the right-hand side uses an `if` statement for control flow and increments the value of variable `s` using an assignment, making it arguably easier to understand. Notwithstanding, confusion is a direct consequence of the experience of the person who is reading the code. A developer who customarily employs the idiom on the left-hand side may find it easier to understand.

Recent work [6, 4, 7, 8] has attempted to elicit and catalog simple code patterns and language constructs that tend to confuse developers when reading them and for which there are less confusing, functionally-equivalent alternatives. These confusing constructs and patterns are called "atoms of confusion" [4] when it is possible to experimentally ascertain that there is a less confusing alternative (otherwise they are called "atom candidates").

```
res = properties.reduce((s,
    b) => {
  b && s++
  return s
}
```

```
res = properties.reduce((s,
    b) => {
  if(b) {
    s = s + 1
  }
  return s
}
```

Listing 1: Code snippet from NERVJS/TARO project (left-hand side) and an alternative implementation (right-hand size).

The snippet on the left-hand side of Figure 1 contains two atoms of confusion discussed in previous work [4] focusing on the C language: Logic as Control Flow and Post-increment. Previous work has identified atoms of confusion in two languages, C and Java. According to the studies of Gopstein et al. [4], removing the atoms Logic as Control Flow and Post-increment from small C code snippets improved the ability of study participants to predict their outcomes by 41% and 34%, respectively. For small Java code snippets, Langhout and Aniche [7] report improvements of 53% and 46.27%, respectively. These results show that atoms of confusion may be **challenging** for developers using these languages. In addition, a study [9] with 14 large-scale open source projects written in C found that 4.38% of the lines of code have an atom and their presence has a strong correlation with bug fixing commits and long code comments. These results show that atoms of confusion are **prevalent** in large systems. They are also **relevant**, even for experienced software developers.

The goal of this paper is to investigate whether atoms of confusion that were identified in Java and C also cause confusion to JavaScript developers. The latter, in spite of the syntactic similarities to Java, is a dynamically-typed language that is arguably more related to Scheme than it is to Java [10]. JavaScript's programming culture differs from that of languages such as C and Java, particularly due to its dynamic capabilities and weaker type system. At the same time, the language is not so different that previously identified atoms cannot be represented in it. We present the results of a mixed-method research effort based on two experiments, a set of interviews, and a mining software repositories study. The two experiments ask the participants to predict the output of small code snippets, where some of them have atom candidates and others do not. These two experiments were de-

3

signed and conducted independently, by different researchers, with different samples, and complementary methodologies; one is uses a repeated measures design [11] whereas the other uses a Latin square design [12]. Both our experiments suggest that two code patterns that have been previously observed to confuse C programmers also confuse JavaScript programmers (with statistical significance, after correction): the comma operator and assignments being used as values. In addition, pre- and post-increment operators, when used as values, omitted curly braces, changes of literal encoding, and type conversions have led to confusion in at least one of the two experiments. For all these cases effect sizes were either medium or high. In addition, we have discovered that object destructuring and automatic semicolon insertion are potential JavaScript-specific atoms that have not been discussed in previous work.

To learn more about the misunderstandings in the snippets used in the experiments, we also interviewed 15 experienced professional developers. Participants preferred the code snippets without the atoms in 70% of the cases. The frequency with which the atoms appear *in the wild* might reveal cases where developers should be more or less careful when writing JavaScript code. To measure the prevalence of the atoms, we mined popular JavaScript projects from GitHub looking for the atom candidates that were investigated in both experiments. We found out that atoms frequently appear in practice. Our MSR effort reveals that atom candidates are frequent and used intensively in 72 popular open-source JavaScript systems.

2. Related Work

The concept of program comprehension is central to software maintenance and feature development [13, 14]. The study of program comprehension in the presence of atoms of confusion was introduced by Gopstein et al. [9], who defined them as small code patterns that can verifiably lead to misunderstandings. In addition, for a code pattern to be considered an atom, there must be some alternative pattern or language construct that is functionally equivalent and less likely to cause confusion. One example that occurs in many programming languages is the Change of Literal Encoding. For instance, the C code statement `printf("\%d", 013);` often leads programmers to predict the output to be `13`, even though the correct answer is `11`. This occurs because a leading `0` in a numerical literal indicates that the number is in base `8`, a fact that is not only unknown to less experienced program-

mers, but misleading even for seasoned developers. After formulating a list of 19 atom candidates, Gopstein et al. [4] were able to identify 15 code idioms and language constructs that present a statistically significant difference in comparative answer correctness when each of the 15 atoms was removed. The least confusing atom produced a 14% boost in prediction accuracy when it was removed, whereas the most confusing one showed a 60% accuracy increase. In a different work, Gopstein et al. [9] presented the results of a comprehensive study on the incidence of atoms of confusion in the wild, considering open-source projects written in C and C++ programming languages.

These previous studies [4, 9] have a great influence in our work. Even though JavaScript is a dynamic language that differs from C in a variety of aspects, it also has a number of constructs in common. In addition, the imperative aspects of the language are syntactically similar to C, e.g., assignments, conditional expressions, `if`-statements, pre and post increments and decrements, among others. On the one hand, this means that some of the atoms of confusion that exist in C programs may also occur in JavaScript code. On the other hand, differences between the languages and the surrounding programming cultures may lead to code patterns that are atoms of confusion in one language not being atoms in the other one. Compared to previous research, in this work we leverage the *Latin square* experimental design to control both participant experience and education better, while investigating the effect of the atom candidates in code comprehension. We present details about our setup in Section 3. Differently, previous research works (e.g., [4, 7]) use a *Randomized Partial Counterbalanced Design*, which does not *block* relevant characteristics of the participants.

Oliveira et al. [8] conducted an experiment using an eye-tracker with 30 participants involving three atoms in C. They found that code with atoms of confusion requires more time from participants to predict the output and more visual effort to comprehend. Gopstein et al., in another paper [15] further explored the original atoms by additionally conducting interviews and having programmers discuss among themselves the studied atoms. The authors argue in favor of complementing a quantitative analysis. Indeed, their study revealed findings such as a *"correct evaluation of an atom might not mean that a programmer understood its meaning"*. Here, we also complement the quantitative analysis with interviews.

Langhout and Aniche [7] derived a set of 14 atoms for Java and performed an experiment with 132 students. For seven atom candidates (out of 14), they

report that participants are more likely to make mistakes in the confusing version of the code. Castor [5] presented a preliminary catalog of six atom candidates for the Swift programming language. Unlike JavaScript, in Swift, most of the atoms identified by Gopstein et al. [4] are avoided by construction, e.g., it does not have assignments as values, increment operators, or macros.

Medeiros et al. [6] analyzed 50 open source projects written in the C language with the goal of evaluating 12 code patterns called "misunderstanding patterns" by the authors. Many of these misunderstanding patterns are either atoms of confusion or atom candidates [4]. This study shows that these patterns are prevalent; among these 50 projects, there are more than 109K occurrences of misunderstanding patterns. In order to gauge the relevance of these misunderstanding patterns, the authors sent 35 pull requests removing occurrences of these patterns to randomly-selected open source projects. The authors of the study received feedback for 21 of these pull requests and the maintainers of the projects accepted 8 of them (22.86%).

3. Methodology

The main goal of this research is to investigate the impact of atoms of confusion on JavaScript code comprehension. As such, in this paper we answer the following research questions:

(RQ1) What is the impact of atoms of confusion on the comprehension of JavaScript code?

(RQ2) Do JavaScript developers identify atoms of confusion as contributing to program misunderstanding?

(RQ3) What is the frequency of occurrence of atoms of confusion in practice (i.e., in open-source JavaScript projects)?

> **(castor note)** This one will probably be removed.

By answering these research questions, we can either generalize or refute the findings about atoms of confusion already discussed in the literature (goal of research questions (RQ1) and (RQ2)).

> **(castor note)** The following pertains to the third RQ.

In addition, answering the third question allows us to enrich existing catalogs about atoms of confusion and discuss how often they occur in practice

(castor note) This should probably be removed.

. Answering these research questions also lays the foundations for the implementation of tools that can automatically transform code into cleaner versions. We conduct a mixed-methods study to answer these questions. It includes two independently designed and conducted experiments: a repeated measures/within subjects study and one using a Latin square, counterbalancing experimental design to control for subject experience. By using two different experiments, designed by different researchers, with different designs, and conducted with different participants, we hope to rigorously verify the impact of atoms of confusion in code comprehension activities. These two experiments explore the impact of atom candidates on understanding JavaScript code (goal of research question RQ1). They investigate whether or not programs that contain atom candidates tend to produce more misunderstanding for programmers trying to predict their outcome.We also conduct a set of interviews with developers to contrast the quantitative results of the experiments with their preferences and opinions (thus addressing RQ2).

3.1. Repeated Measures Study

(castor note) We must have introduced the atoms before.

The first experiment we perform uses a repeated measures, or within-subject, design [11]. This experiment is inspired by the work of Gopstein et al. [4], but targeting JavaScript instead of the C language, and including some methodological differences. It covers all the atom candidates analyzed in the work of Gopstein et al. that can be transposed to JavaScript and adds candidates that are specific to this language.

Experimental design

In this experiment, the control variable consists of a tiny program containing a single atom candidate and the treatment is a functionally-equivalent version that does not contain the atom. The main dependent variable is whether subjects are able to correctly determine the output of the programs. The other dependent variable is the time required to correctly vs. incorrectly determine the outputs, independently of the presence of atoms. We use the same null hypothesis as Gopstein et al.: *"code from both control and treatment groups can be hand-evaluated with equal accuracy"* and the alternate

hypothesis that *"the existence of an atom of confusion causes more errors than other code in hand-evaluated outputs."*. For each atom candidate, we built six programs, three versions including the atom candidate and three versions not including it.

To identify atom candidates, we start out by selecting every atom candidate from Gopstein's study that can also happen in JavaScript. That excludes atoms related to pointers and preprocessor directives. Furthermore, we conduct an informal search of programming forums, solutions to Code Golf [1] problems, code style guidelines, and expert knowledge to identify new atom candidates. Overall, the repeated measures study analyzes 24 atoms, 15 from the original study and 9 that are new. Some of these atoms are also analyzed in the Latin square study (Section 3.2). Table 1 presents the atom candidates that both studies investigate and Table 2 presents the ones that only the repeated measures study analyzes.

For each atom candidate, we write six short programs, three including and three not including the atom candidate. The former are the control for the study, or the **obfuscated** versions of the programs. The alternative, atom-free, versions of the programs are the treatment, or **clean**, versions. Each subject is exposed to one obfuscated and one clean version of each atom, totalling 48 programs, and should determine their outputs. As mentioned before, we measure answer correctness and the time to answer questions correctly vs. incorrectly. We control for learning effects [17] in three manners. First, by having multiple obfuscated and clean versions for each atom and presenting only one of each per subject. Second, by presenting the programs in a random order. Third, by presenting obfuscated and clean versions corresponding to the same atom candidate with at least 11 other programs in between, in accordance to the original experimental protocol [4].

Study instrument

The experiment is conducted by means of a questionnaire made available as a web application. As part of this effort, we carried out an informal pilot whose main objectives were (i) to spot bugs in the application and in the data collection mechanism; (ii) to gain feedback from respondents about the user experience of the application; and (iii) to formulate an estimate about how long answering the questionnaire would take on average. Based on the

---

[1]https://codegolf.stackexchange.com/

Table 1: Atom candidates analyzed in the two experiments. Nine of these atoms come from the original study of Gopstein and colleagues [4]. For these atoms, both the descriptions and the examples, with the required adaptations, are taken directly from their paper.

Atom Name	Description	Obfuscated	Clean
Logic as Control Flow	Traditionally, the && and —— operators are used for logical conjunction and disjunction, respectively. Due to short-circuiting, they can also be used for conditional execution.	V1 && F2();	if (V1) { F2() };
Assignment as Value	The assignment expression changes the state of the program when it executes, however it also returns a value. When reading an assignment expression people may forget one of the two effects of the expression	V1 = V2 = 3	V2 = 3 / V1 = V2
Pre-Increment / Decrement	Similar to post-increment/decrement, these operators change a variable's value by one. In contrast to the other operators, pre-increment/decrement first update the variable then return the new value.	V1 = ++V2	V2 += 1 / V1 = V2
Post-Increment / Decrement	The post-increment and decrement operators change the value of their operands by 1 and return the original value. Confusion arises because the value of the expression is different from the value of the variable.	V1 = V2++	V1 = V2 / V2 += 1
Comma Operator	The comma operator is used to sequence series of computations. Whether due to its eccentricity, or its odd precedence, the comma operator is commonly misinterpreted.	V3 = (V1 += 1; V1)	V1 += 1; / V3 = V1;
Implicit Predicate	The semantics of a predicate are easily mistaken. The most common example happens when assignment is used inside a predicate or when a success state is represented as 0.	if (4 % 2)	if (4 % 2 != 0)
Conditional Operator	The conditional operator is the only ternary operator in C, and functions similarly to an if/else block. However, it is an expression for which the value is that of the executed branch.	V2 = (V1==3)?2:V2	if (V1 == 3) { / V2 = 2; / }
Automatic Semicolon Insertion	Semicolons in JavaScript are optional in some contexts. However, some cases where semicolons are optional can confuse the JavaScript parser. This may lead to surprising behavior.	let V1 = 9 / (console.log(V1))	let V1 = 9; / (console.log(V1));
Omitted Curly Braces	When control statements omit curly braces, the scope of their influence can be difficult to discern	if (V) F(); G();	if (V) {F();} G();
Arithmetic as Logic	Arithmetic operators are capable of mimicking any predicate formulated with logical operators. Arithmetic, however, implies a non-boolean range, which may be confusing to a reader	(V1-3) * (V2-4)	V1!=3 && V2!=4

Table 2: Atom candidates analyzed only in the repeated measures experiment. Six of these atoms come from the original study of Gopstein and colleagues [4]. For these atoms, both the descriptions and the examples, with the required adaptations, are taken directly from their paper.

Atom Name	Description	Obfuscated	Clean
Type Conversion	JavaScript will often implicitly convert types when there is a mismatch.	{} + 4	{}.toString() + (4).toString()
Repurposed Variables	Sometimes this conversion also results in a different outcome than a reader expects. When a variable is used in different roles in a program, its current meaning can be difficult to follow	let n = 1; n = "OK";	let n = 1; let s = "OK";
Change of Literal Encoding	All numbers are stored as binary, but for convenience we represent numbers in decimal, hex, or octal. Depending on the circumstance, certain representations are more understandable.	let V1 = 013 console.log(V1)	let V1 = 11 console.log(V1)
Infix Operator Precedence	JavaScript has more than 40 operators[a], the majority of them infix, each in one of 17 precedence levels with either right or left associativity. Most programmers know only a functional subset of these rules.	0 && 1 —— 2	(0 && 1) —— 2
Constant Variables	Constant variables are a layer of abstraction that emphasize a concept rather than a particular value itself. When trying to hand evaluate a piece of code having a layer of indirection can obscure the value of the data.	let V1 = 5 console.log(V1)	console.log(5)
Dead, Unreachable, Repeated	Redundant code is executed to no functional effect, but its appearance may imply that meaningful changes are being made	V1 = 1; V1 = 2;	V1 = 2;
Lack of Indentation, no Braces	Block delimiters and indentation are very important to understand nested structures in programs. If block delimiters are missing and indentation is incorrect, developers may misunderstand the nested structure of the program.	if (V1) if (V2) V3 = V3 + 2; else V3 = V3 - 1;	if (V1) { if (V2) { V3 = V3 + 2; } else { V3 = V3 - 1; } }
Lack of Indentation, with Braces	Similar to Lack of Indentation, no Braces, but the block delimiters are present. The lack of indentation still has the potential to cause confusion.	if (V1) {}; V1 = V1 - 1;	if (V1) {}; V1 = V1 - 1;
Property Access	Object properties can be accessed using the dot (".") notation used in most object-oriented languages and using notation similar to how elements in collections like dictionaries and maps are accessed in other languages. The latter is less popular[16] than the former. Lack of familiarity may lead to confusion.	let V1 = {P1:0, P2:5 } V1["P1"] = 9;	let V2 = {P3:0, P4:5 } V2[P3] = 9;
Arrow Function	The usage of arrow function notation, instead of the more traditional JavaScript notations for anonymous functions.	let V1 = x =>x * x * x; let V2 = V1(5);	let V3 = (x) =>{ return x * x * x; } let V4 = V3(5);
Array Spread	Spread facilitates copying, extracting elements and concatenating arrays. Some caveats exist as the copying is shallow so this can cause misunderstanding between programmers.	let V1 = [1,2]; let V2 = [5]; V3 = [...V1, ...V2]	let V1 = [1,2]; let V2 = [5]; V3 = [V1[0], V1[1], V2[0]]
Object Spread	Similar to the spread in arrays, but with slightly more complex semantics because it may lead to the values of properties being overwritten if properties with the same names exist in two objects combined by means of spread.	let V1 = {P1:5, P2:8, P3:13}; let V2 = {...V1, P3:99};	let V1 = {P1:5, P2:8, P3:13}; let V2 = {P3:99}; V2.P1 = V1.P1 V2.P2 = V1.P2
Array Destructuring	This feature provides an approach to perform a limited form of pattern matching based on the elements of an array. The syntax mixes the introduction of new variables with array literals and is arguably less explicit than directly initializing new variables with array elements.	let [V2,,V4] = [1,2,3,4]; console.log(V2, V4);	let V1 = [1,2,3,4]; console.log(V1[0], V1[3]);
Object Destructuring	Similar to array destructuring, but uses the syntax associated with object literals to introduce new variables by matching their names with the names of properties of the destructured object. This is less explicit than initializing the new variables with the values of said properties.	let V1 = {P1: 5}; let {P1, P2 = 20} = V1;	let V1 = {P1: 5}; let P1=V1.P1, P2=20;

[a] https://developer.mozilla.org/en-US/docs/Web/JavaScript/Guide/Expressions_and_Operators

feedback of the participants of the pilot, we present only one obfuscated and one clean version of each atom candidate, to reduce to total time of the study and potentially increase participation.

We organize the study in three sections. In the first one we present instructions and also a check button whose checking means users agree that all collected data will be used anonymously and solely for research purposes. The instructions explain how the study works and asks them to dedicate their attention to it. We stress to participants the importance of not using any aids during the experiment, such as online or console interpreters.

The second section presents the programs, one at a time. For each one, there is a text box where the answer should be written and a "Next" button. If the subject leaves the text box empty and presses "Next", this is treated as a wrong answer. To reduce the likelihood of a subject attempting to execute the programs, the web application verifies whether the subject tries to change tabs or windows and presents a pop-up message if that is the case. Furthermore, it presents the programs as images, instead of text, to demotivate respondents from resorting to external resources by copying and pasting the code into an interpreter. Upon submitting an answer for a particular program, the subject is automatically led to a similar page, containing the next program.

We do not provide feedback about the time subjects take to answer each question. We also do not tell them whether their answers are correct or not. This aims at avoiding introducing bias for future respondents. Since we posted information about the study on social media platforms, possible threats could have arisen if we gave respondents instant feedback.

Study audience

We posted the link to the questionnaire on the authors' social media platforms, on Brazilian CS departments' mailing lists, and on two programming subreddits. We explained our research purposes and asked developers to take part in the study only if they have some familiarity with JavaScript. We collected 70 responses.

3.2. Latin Square Study

This experiment uses a different experimental design that makes it possible to control for the experience of the subjects. It employs a subset of the atom candidates used in the repeated measures study. More specifically, this study investigates ten of the atoms investigated in the repeated measures

study: nine previously-validated atoms of confusion for the C language [4] that also exist in JavaScript programs plus one atom candidate that is specific to the JavaScript language (Automatic Semicolon Insertion).

Experimental design

The design of the Latin square study blocks two variables (subject experience and the programs) and considers two treatments: the presence or absence of atom candidates within the programs. To achieve such a design goal, controlling the effect of experience and individual programs, we resort to the *Latin Square Design* [12]. Using this design we create a 2 x 2 matrix in which each row represents a subject and each column indicates the set of programs. The design of each square (a replica) is such that no treatment is repeated in the same row or column. For example, considering that we have a set of 10 code samples $s_0, s_1, ..., s_{10}$, if a given subject (P1) is asked to predict the output of the code samples $s_0, s_1, ..., s_5$ that contain atom candidates, then, when answering questions about clean programs, they will only be presented with clean versions of the code samples $s_6, s_7, ..., s_{10}$. Furthermore, a given subject (P2), who constitutes the second row of our example square, will be asked questions about the clean versions for $s_1, s_2, ..., s_5$, and will answer questions about obfuscated programs for $s_6, s_7, ..., s_{10}$. By doing that, we guarantee that all versions of the programs are contained within each square, and that each configuration occurs only once within a square. Figure 1 offers a visual representation of the concept.

Figure 1: Latin square design. Each "square" corresponds to a replica in our study. Each replica comprises two participants (square rows, e.g., P1 and P2) and two sets of programs (CS-1 and CS-2). We randomly apply the treatments (atom or non-atom code) to the cells of the squares.

For each of the 10 selected atom candidates, we write one short program containing the atom. As in the repeated measures study, we call them the obfuscated versions of the programs. We also write corresponding short programs without the atom candidates and call them the clean versions of the programs. Overall, the experiment employed 20 programs, 10 obfuscated and

10 clean. In order to reduce the cognitive effort, each subject was asked to predict the output of 10 programs, five obfuscated and five clean. The order in which they are presented is randomized. By doing this, we seek to minimize the chances of subjects being aware that the current listing they are analyzing contains (or not) atoms of confusion. That is, each subject should indicate what would be the outcomes of the programs, some of them having atoms of confusion (while other programs do not). Since participants are not exposed to obfuscated and clean versions corresponding to the same atom candidate, learning effect is not possible. We measure answer correctness and the total time each participant requires to participate in the experiment.

Study instrument

We implemented our experiment as a questionnaire in a web application and carried out an informal pilot.Undergraduate students and professional colleagues took part in the pilot. Some users reported layout defects, and many reported that the landing page did not explain the study well enough. We also spotted minor issues with our routines to create and populate the latin squares.

We organize the questionnaire in three sections. The first section aims to characterize the subjects, asking their age, education level, and programming experience. We also include a check button, whose checking means users agree that all collected data will be used solely for research purposes. In the second section, we present to the participants a small set of instructions, where we explain how the study works and ask them to dedicate their attention to it. We stress to participants the importance of not using any aids during the experiment, such as online or console interpreters. For each question page, we kept track of whether or not the subjects switch windows.

The last section of the questionnaire presents a sequence of ten questions, each containing a program. For each question, there is a text box where the answer should be written. There is also an "I do not know" button, which, when clicked, leads the subject to the next question. In our setting, "I do not know" is treated as a wrong answer. The programs are also presented as images copied from a text editor. Upon submitting their answer for a particular question, the subjects are automatically led to a similar page, containing another program. Similarly to the repeated measures study, we do not provide feedback about time and correctness to the subjects.

Study audience

We posted the invitation to take part in the study on a JavaScript Reddit channel.[2] We explained our research purposes and asked developers of any level of expertise to take part. Within twelve hours, we collected more than 150 answers, populating more than 70 replicas of the Latin Squares. We collected significant data on time taken and discrepancies in answer correctness between obfuscated and clean versions of the programs. Some inconsistencies arose while building the squares, for instance, when a user quit in the middle of the questionnaire. We discarded from our study all squares that contained incomplete rows (a common approach for data imputation).

3.3. Interview Study Setting

To complement the two experiments, we perform semi-structured interviews with professional JavaScript developers, aiming to identify their perceptions regarding programs containing atoms of confusion (research question RQ2). We also ask each participant if they know of any other JavaScript-specific construct or idiom they think is likely to make the code hard to understand. In this section we detail the protocol we followed to conduct the interviews and to analyze the results.

Participant Selection

We invite the participants of the interviews using a snowballing technique. That is, starting from our network of contacts, we invite an initial set of candidates to take part in our experiments. From this initial list, we ask for an indication of additional candidates. Our main selection criterion is that all participants should have been working with JavaScript in their daily professional activities. We invited a total of 17 developers, and 15 of them agreed to participate.

Interview Process

We conduct semi-structured interviews using web conferencing software. We record all the interviews with the consent of the participants. On average, the interviews last 26.29 minutes, with the shortest one lasting 14.59 minutes and the longest one 43.06 minutes. Two of the researchers conduct the interviews, and a third one listens to all the recordings to cross-validate the

---

[2]https://www.reddit.com/r/javascript/

collected data. The interviews have three main parts. In the first one, we ask the developers the following demographic information: name, email, gender, level of education, current job position, JavaScript experience in years, and other programming languages they have worked with. Table 3 summarizes this demographic information.

Table 3: Demographic information of the participants

ID	Education	JS Experience	Other Languages
P1	BSc Degree	9 years	Java, PHP, C, Go
P2	HS Degree	3 years	Python, Go, Dart, Lua, C++, C#
P3	BSc Degree	4 years	Java, C
P4	Undergraduate	3 years	Python, C, C++, Java, Go
P5	Master Student	3 years	Python, SQL
P6	Bsc Degree	15 years	Java, PHP, C, Python , Ruby, C#
P7	BSc Degree	6 years	C, C++, Java, Assembly, Kotlin
P8	PhD Degree	2 years	Java, Python
P9	BSc Degree	5 years	PHP
P10	Master Student	4 years	Java, C# and Python
P11	Master Student	4 years	Java, Erlang, C#, Cobol
P12	Master Degree	1 year	C
P13	Master Degree	13 years	Java, PHP
P14	Master Student	3 years	C, Python, Ruby
P15	BSc degree	2 years	Java, PHP

In the second part of the interview our aim is to allow the subjects to describe their JavaScript experience, as well as to allow them to reveal any JavaScript constructs they regard as innately confusing. This allows us to identify potential atom candidates that are more specific to the JavaScript language. Examples of questions we explore in this section include: *Does JavaScript favor developers to produce code that is hard to understand?* and *Do you regard any particular construct or idiom of the language as especially confusing?*

In the third part of the interview, participants are shown pairs of programs that were used in the experiments, where each pair contains a clean and an obfuscated version of the same atom candidate. For this part we only use atoms that appear in both experiments. The participants are asked to evaluate which version of the code is easier to understand. To avoid introducing bias in the answers, the interviewers do not explain that one of the versions in each pair contains the atom under investigation. Subjects are just presented the pairs of programs and allowed to take the necessary time to decide on the most readable one.

Interview Analysis

We first transcribe each interview and then examine the broad distribution of the answers. Our goal is to build an initial understanding of the participants' perceptions with regards to the challenges to understand JavaScript code in general and JavaScript code with atoms of confusion in particular. We follow-up with an open-coding procedure, highlighting the main themes and quoting the answers of the participants. We present these results in Section 6.

3.4. Mining JavaScript Software Repositories Setting

To understand how often the analyzed atom candidates appear in open source projects, and thus answer our third research question (*What is the frequency of occurrence of atoms of confusion in practice?*), we mine a set of GitHub open source repositories. To this end, we first collect the most popular GitHub repositories that are primarily written in JavaScript. We measure popularity using the project's stargazers. This metric, available through the GitHub API, represents the number of stars a project received from users of the platform. The same metric has been used in a number of previous studies as a proxy to estimate a project's popularity [18, 19]. We then select the top 100 most popular repositories and remove projects that do not reach the first quartile of the distribution of lines of code.

After filtering out JavaScript project candidates, in the second step we build a curated dataset comprising the top 72 repositories. Examples of projects in this dataset include REACT, NODE JS, and ANGULARJS. Table 4 presents some statistics about the projects we consider in our research. The size of the projects range from small ones (5543 lines of code) to complex systems with more than 1 MLOC. All projects in our dataset have at least 1244 forks and at least 23672 stars. We automate all the steps to filter, clone, and collect the statistics from the repositories using Python scripts.

We mine the occurrence of atom candidates from the repositories in our curated dataset using source code queries that we wrote using the CodeQL language [20]. CodeQL is an object-oriented variant of the Datalog language [21], and currently supports researchers and practitioners to query the source code of systems written in different languages (such as C++, Java, and JavaScript). We also automate the process of running the queries and exporting the results to a format that simplifies our analysis (and also the reproduction of this study). Finally, we compute some descriptive statistics to measure the prevalence of atoms of confusion in practice.

Table 4: Some descriptive statistics about the projects used in the MSR study

	Min.	Median	Mean	Max.
Lines of Code	5543	36161.5	111432.33	1278405
Num. of Forks	1244	6078	8906.24	68849
Num. of Contributors	6	285.5	533.44	4047
Num. of Stars	23672	34990	46919.51	310935

4. Results of the Repeated Measures Study

In this experiment, we investigate whether 24 atom candidates impact readability. We examine how the performance of the participants differ when predicting the outputs of the clean and obfuscated versions of small code snippets (Section 4.1). Based on those results we establish which candidates can actually be considered atoms of confusion. In addition, similarly to Gopstein et al. [4], we verify whether there is a trade-off between correctness and time irrespective of the presence of atoms; we check if participants who correctly predict the outputs of the programs tend to do so more slowly than the ones that make incorrect predictions (Section 4.2).

This study had 70 participants. Out of these, 67 have received university-level training in programming, 4 hold a PhD degree, 7 hold a master's degree, and 11 hold a bachelor's degree. The median programming experience of the participants is four years (mean 6.5 years) and the median experience with JavaScript programming is one year (mean 2.5 years).

4.1. Correctness Analysis
Exploratory Data Analysis

For each of the 24 atom candidates, the participants predicted the outputs of two versions, one clean (control) and one obfuscated (treatment), as per the definition of a repeated measures design [12]. Thus, each participant examined 48 code snippets. We also measured the time required by the participants to provide their answers. For most cases, participants made more mistakes when attempting to predict the outcomes of the obfuscated versions.

Table 5 presents the number of correct predictions for the clean and obfuscated versions of each atom candidate, as well as the percentage difference. In only three cases the participants predicted results correctly more often for the obfuscated versions, Arithmetic as Logic, Logic as Control Flow, and Property Access. The former two were empirically validated as atoms of confusion

17

Table 5: Summary of the correctness analysis.

Atom	Obfuscated	Clean	$\Delta(\%)$
Type Conversion	5	27	+440
Change of Literal Encoding	9	36	+300
Comma Operator	8	29	+262
Object Destructuring	15	42	+180
Assignment As Value	18	43	+139
Pre-Increment	13	28	+115
Post-Increment	8	17	+112
Array Destructuring	15	28	+87
Indentation no Braces	30	50	+67
Omitted Curly Braces	35	51	+46
Object Spread	21	30	+43
Array Spread	22	31	+41
Automatic Semicolon Insertion	15	21	+40
Infix Operator Precedence	47	53	+13
Arrow Function	54	58	+7
Implicit Predicate	41	44	+7
Indentation With Braces	53	55	+4
Conditional Operator	60	62	+3
Constant Variables	59	61	+3
Dead Unreachable Repeated	63	64	+2
Repurposed Variables	23	23	0
Arithmetic as Logic	52	47	-10
Logic as Control Flow	33	28	-15
Property Access	51	42	-18

in the study of Gopstein and colleagues [4]. For seven of the atom candidates, the number of correct predictions for the clean versions is more than twice the corresponding number for the obfuscated versions. Some of these atoms are common programming patterns in JavaScript programs, such as pre- and post-increments and object destructuring. The difference was even greater for Type Conversion (+440%) and Change of Literal Encoding (+300%).

Statistical analysis

Since this study uses a within-subject design, the data is paired and dependent. Our goal is to verify whether there is a significant difference in the performance of the same participants when examining clean and obfuscated versions of programs. The analyzed data is also dichotomous, since each participant either correctly predicts the output of a program or does not. Considering this scenario, we employ McNemar's test [22], which is aimed at *"judging the significance of the difference between correlated proportions"*.

The "McNemar test" column of Table 6 indicates the p-values for the tests. These p-values are compared against an alpha of 0.0021, after applying Bonferroni correction (0.05/24). One of the atom candidates, Constant

Variables, is omitted from the table because, due to the very low number of mistakes, it is not possible to perform the test. The table shows that eight of the 24 analyzed atom candidates can be considered atoms of confusion based on the results of our study, Object Destructuring, Type Conversion, Change of Literal Encoding, Comma Operator, Indentation No Braces, Assignment as Value, Pre-Increment, and Ommitted Curly Braces. Among these eight, only Object Destructuring is JavaScript-specific. Six of the other atoms have been previously reported for the C language [4] and five for Java [7]. Indentation No Braces is not specific to JavaScript, but has not been studied previously. For some candidates that have been previously identified as atoms, there was not a statistically significant difference. This is the case for Repurposed Variables, Logic as Control Flow, Implicit Predicate, and Conditional Operator.

We leverage the odds ratio (OR) as a measure of effect size, to quantify the meaningfulness [23] of these results. The odds ratio is an appropriate measure of effect size for scenarios where McNemar's test is applicable [24]. For example, the odds ratio for Type Conversion is 12. This suggests that the odds of a participant correctly predicting the output of a clean version while incorrectly predicting the output of the corresponding obfuscated version is 12 times higher than the odds of correctly predicting the output of the obfuscated version and incorrectly predicting the output of the correct one.

According to the thresholds established by Chen and colleagues [25] (OR = 1.68 small, 3.47 medium, and 6.71 large), for most of the atoms where there is a statistically significant difference the odds ratio can be considered high. The value of the odds ratio for Object Destructuring is ∞ because no participant correctly guessed the output for the obfuscated version while incorrectly guessing the output for the clean one, which results in a denominator of 0. Omitted Curly Braces exhibits an OR of 6.3, which is borderline between medium and high [25]. Array Destructuring and Post-Increment are two atom candidates that deserve a special note. The difference in the performance of participants examining clean and obfuscated versions of these atoms was not statistically significant, after correction. Notwithstanding, effect sizes were large and medium, respectively. This is a result that hints at practical relevance in spite of the p-values [23]. Further investigation is required for these two atom candidates.

Table 6: Hypotheses Testing (correctness). Asterisks (*) indicate a statistically significant difference. Using Bonferroni correction, a p-value is considered statistically significant if it is lower than 0.0021.

Atom candidate	McNemar test	odds ratio
Object Destructuring	**< 0.0001***	∞
Type Conversion	**< 0.0001***	12.0
Change of Literal Encoding	**< 0.0001***	28.0
Comma Operator	**< 0.0001***	11.5
Indentation No Braces	**< 0.0001***	11.0
Assignment As Value	**< 0.0001***	9.3
Pre-Increment	**0.0003***	16.0
Omitted Curly Braces	**0.0009***	6.3
Array Destructuring	0.0023	7.5
Post-Increment	0.0225	5.5
Array Spread	0.0490	3.3
Object Spread	0.0636	2.8
Property Access	0.0636	2.8
Infix Operator Precedence	0.1796	2.5
Arrow Function	0.2891	3.0
Arithmetic As Logic	0.3323	1.8
Automated Semicolon Insertion	0.3915	1.4
Logic As Control Flow	0.4421	1.5
Implicit Predicate	0.6900	1.3
Indentation With Braces	0.7266	1.7
Conditional Operator	0.7905	1.3
Dead Unreachable Repeated	1.0000	1.5
Repurposed Variables	1.0000	1.0

4.2. Time Analysis

Exploratory Data Analysis

In this section we investigate whether there is a trade-off between correctness and time required to predict program outputs, similarly to the analysis performed by Gopstein et al. [4]. For this analysis, we consider only whether participants correctly predict the outputs of programs or not, independently of the presence or absence of atoms, and the time required to do so. The null hypothesis we investigate is that there is no significant difference in time between correct and incorrect predictions. Intuitively, we would expect correct responses to take longer, on the basis that participants would spend more time carefully analyzing the programs before providing an answer.

Table 7 presents the mean and median times required by the participants to provide answers. The data is based on the 1648 correct and 1572 incorrect answers collected in the study. Contrary to our intuition, the mean time required by participants to provide an incorrect answer is 25% greater than the time required to provide a correct one and the median time is 20% greater.

Table 7: Time in seconds required by the participants to correctly and incorrectly predict the outcomes of the programs.

	Mean	Median	Std. Dev.
Correct	35.44	24.36	38.43
Incorrect	44.50	29.29	62.48

The standard deviations suggest that there is great variation in the times required by different participants, for different code snippets. Notwithstanding, the standard deviation for incorrect answers is 62.58% greater when compared to the same statistic for correct answers.

Statistical Analysis

We employ the non-parametric Mann-Whitney U test, also known as Wilcoxon's rank-sum test, to verify the aforementioned null hypothesis, i.e., whether the times for correct and incorrect answers can be considered identical. That test produces a p-value of **0.00000046**, which shows that the difference is statistically significant and we can reject the null hypothesis. To gauge the magnitude of that difference we employ Cliff's Delta as a measure of effect size. This produces a result of **0.1027**, which suggests that, although participants tend to provide correct predictions more quickly, they do that only slightly so.

5. Results of the Latin-Square Study

In the latin-square study, we estimate the impact of atoms considering the same two perspectives of the previous experiment: *correctness* (number of wrong answers) and *time* (how long to provide a correct answer). In particular, in the time analysis, differently from the repeated measures study, we compare the time required by participants to submit a correct response when analyzing code snippets with and without atom candidates. This study is based on the responses of 140 participants (a total of 70 replicas). All participants had taken at least some university course or hold a bachelor degree or equivalent. In addition, 21 participants hold a master's degree and three a doctorate degree. Considering the programming experience of our respondents, 19% have more than ten years of programming experience, 37% have between four and ten years of experience, 37% have between one and four years of experience, and 7% have less than one year of programming

experience. Accordingly, we characterize the effect of atoms of confusion in JavaScript code taking into account the perceptions of both novice and experienced developers.

5.1. Correctness Analysis

Exploratory Data Analysis

As mentioned in Section 3.2, each participant in this experiment evaluated ten code snippets, from which five were in their obfuscated versions, whilst the other five contained clean versions of the code snippets (i.e., without the atom candidates). As in the repeated measures study, the participants should provide the expected outcomes of the code snippets. Differently from that study, in this one each participant only predicts the outcome of one version of each code snippet, either clean or obfuscated. As discussed, we collect information about *correctness* (whether the participant correctly predicted the program's output) and *time*. Table 8 and Figure 2 summarize the results of the correctness experiment.

Considering Table 8, the clean versions of six code snippets present at least a 15% improvement in answer correctness when compared with the obfuscated versions. In particular, the presence of the *Comma Operator* atom exhibits the highest impact on misunderstanding. Frequently used constructs and idioms, such as *Post Increment* and *Omitted Curly Braces* (see Section 7), also result in many mistakes. The boxplot of Figure 2 shows a non-negligible decrease in the average number of incorrect answers when observing the clean versions of the code snippets. Also, the sample of responses with no atoms had almost no dispersion, which supports the argument that the clean versions of the code snippets are easier to evaluate correctly.

Table 8: Summary of the correctness analysis

Atom	Obfuscated	Clean	$\Delta(\%)$
Comma Operator	28	65	+132
Automatic Semicolon Insertion	32	68	+112
Post-Increment	48	64	+33
Omitted Curly Braces	47	58	+23
Assignment as Value	56	68	+21
Implicit Predicate	58	68	+17
Logic as Control Flow	41	48	+17
Conditional Operator	60	66	+10
Pre-Increment	50	53	+6
Arithmetic as Logic	64	63	-2

22

Statistical analysis

We first use the *Pearson's Chi-squared test* to investigate if there is a statistically significant difference in the frequency of correct and incorrect answers—due to the versions of the code snippets (obfuscated and clean code). The p-values for these tests are reported in the "Chi-square test" column of Table 10. The results indicate that for five atom candidates (Comma Operator, Automatic Semicolon Insertion, Post-Increment, Assignment as Value, Implicit Predicate) the obfuscated versions of the code snippets have a negative impact on code understanding (p-value < 0.05). This result holds even after applying the Benjamini-Hochberg correction with a false discovery rate of 5%.

We measure the effect size of the clean version of the code snippets into the answers' correctness using the *Odds Ratio* (OR). The results are reported in the "Odds Ratio" column of Table 10. In the table, we also report the confidence interval (CI) for the OR in the "CI" column. Although many of the intervals are wide, for six of the atom candidates the lower bound of the confidence interval is greater than or equal to 1. This indicates that, at a 95% confidence level, a developer is likely to commit less mistakes when using the clean versions. For four of the atoms, Comma Operator, Automatic Semicolon Insertion, Assignment as Value, and Implicit Predicate, effect size can be considered large. For Implicit Predicate, it is medium. For instance, when


Figure 2: Number of wrong answers of each subject.

23

comparing clean and obfuscated code snippets pertaining to the Comma Operator atom candidate, we observed an *Odds Ratio* of 19.02. This means that the odds of a correct answer are 19.02 times higher when interpreting the clean version—in comparison with the corresponding obfuscated version of the code snippet. Furthermore, with a 95% confidence level, the odds ratio is between 6.6 and 68.22, i.e., although the error margin is wide, it is strongly in favor of the clean version. If there is a significant likelihood of a developer committing a mistake when analyzing a clean version, the lower bound of the CI should be a number between 0 and 1 (since the OR is a ratio). This is the case for the four atom candidates in the lower part of the table.

Finally, we also use a *Binomial Generalized Logistic Regression* analysis to investigate if either *participant education* or *participant experience* impacts correctness. The findings suggest that *participant experience* impacts the results related to correctness (Figure 3). Even though the presence of atoms of confusion reduces the number of correct answers for all *experience groups*, this effect is not uniform. For instance, novice developers (under one year of experience) provide a higher number of wrong answers for the obfuscated version of the code (57.8%), while developers with more than four years of experience provide more than 70% of correct answers even when evaluating the obfuscated version of the code snippets. It is also important to note that, independently of the *experience group*, developers tend to provide more than 84% of correct answers while predicting the outputs for the clean versions of the code. We replicate the *Pearson's Chi-squared test* for each experience group and find that the statistical difference is less significant for those developers with more than ten years of experience. Differently from *participant experience*, the factor *participant education* does not significantly change the correctness of the answers.

> **(castor note)** I'd like to have the numbers for the results above. What are the p-values, regression coefficients, etc.?

5.2. Time Analysis

Exploratory Data Analysis

Table 9 shows the average time necessary for the participants to give a correct answer about the expected outcomes of a code snippet, considering both obfuscated and clean versions. We do not consider cases where the participants made mistakes. Seven atom candidates required more time from participants to predict a correct response. For these seven atom candidates,

	Clean Code	Confuse Code		Clean Code	Confuse Code
Correct Answer	0.844	0.422		0.869	0.661
Wrong Answer	0.156	0.578		0.131	0.339
	Under one year of experience			One year and under four years of experience	
Correct Answer	0.904	0.723		0.903	0.778
Wrong Answer	0.096	0.277		0.097	0.222
	Four years and under ten years of experienc			More than ten years of experience	

Figure 3: Impact of atoms of confusion on the correctness (over the four participant experience groups)

developers take at least 5.90% less time on average to find a correct answer when considering the clean version of a code snippet. In the extreme case (atom candidate Comma Operator), the participants take 76.23% less time on average to find the correct answer for the clean version of the code snippets. Not all atom candidates, though, require less time for the participants to predict the answer. In fact, for Arithmetic as Logic and Pre-Increment, the time to give a correct answer for the clean versions was 29.06% and 38.19% than for the obfuscated versions—we did not confirm these atom candidates as atoms of confusion in the previous section.

Table 9: Time in seconds to submit a correct answer

Atom	Obfuscated	Clean	$\Delta(\%)$
Comma Operator	87.67	20.84	-76.23
Automatic Semicolon Insertion	46.08	22.04	-52.17
Post-Increment	28.70	25.67	-10.56
Omitted Curly Braces	48.85	30.00	-38.58
Assignment as Value	52.47	48.95	-6.71
Implicit Predicate	36.24	24.01	-33.75
Logic as Control Flow	108.94	51.07	-53.12
Conditional Operator	41.80	39.34	-5.90
Pre-Increment	30.71	42.45	38.19
Arithmetic as Logic	28.82	37.20	29.06

Statistical analysis

We use the *Mann-Whitney U test* to investigate the null hypothesis that developers spend the same amount of time to correctly predict the outcome of clean and obfuscated versions of a code snippet. The results of Table 10 suggest that we should refute the null hypotheses for five atom candidates (Comma Operator, Implicit Predicate, Logic as Control Flow, Pre-Increment, and Arithmetic as Logic), after applying the Benjamini-Hochberg correction

Table 10: Hypotheses Testing (misunderstanding and time). Asterisks (*) indicate a statistically significant difference.

Atom	Chi-square test	Odds Ratio	CI CI	Mann-Whitey U test	Cliff's Delta
Comma Operator	< **0.0001***	19.02	(6.60, 68.22)	< **0.0001***	-0.53
Automatic Semicolon Insertion	< **0.0001***	39.33	(9.21, 356.62)	0.0980	-0.16
Post-Increment	**0.0015***	4.83	(1.73, 15.74)	0.1201	0.15
Omitted Curly Braces	0.0510	2.35	(1.00, 5.77)	0.9253	-0.01
Assignment as Value	**0.0035***	8.39	(1.81, 79.12)	0.8268	0.02
Implicit Predicate	**0.0112***	6.95	(1.46, 66.56)	**0.0029***	-0.29
Logic as Control Flow	0.2920	1.54	(0.73, 3.28)	< **0.0001***	-0.39
Conditional Operator	0.1590	2.73	(0.74, 12.57)	0.2407	-0.12
Pre-Increment	0.7015	1.25	(0.55, 2.85)	**0.0062***	0.27
Arithmetic as Logic	1	0.84	(0.22, 3.12)	**0.0172***	0.23

with a false discovery rate of 5%. For the first three, the analysis suggests that the participants need more time to predict the outcome of the code snippet in the obfuscated code version. For the atom candidates Arithmetic as Logic and Pre-Increment, the participants take less time to predict the outcome of the code snippets in the obfuscated version. We also computed the effect size using Cliff's Delta statistic. In the table, negative effect sizes suggest that it took less time to correctly predict the output of clean versions of the snippets. Considering the thresholds established by Romano et al. [26] ($0.147 < |d| < 0.33$ small, $|d| < 0.474$ medium, otherwise large), we found a large effect for the Comma Operator atom candidate; a medium effect size for Logic as Control Flow, and small effect sizes for Automatic Semicolon Insertion, Implicit Predicate, Arithmetic as Logic, Post-, and Pre-Increment.

6. Results of the Interview Study

In this section we present the results of the interviews with the practitioners. We contrast the findings of the experiments of Sections 4 and 5 with the opinion of software developers about their preferences regarding the obfuscated or clean versions of the code snippets. For the interview study, we only considered atoms that were analyzed in both the repeated measures (Section 4) and the latin square (Section 5) studies.

A total of 15 practitioners took part of the interviews. We collected information regarding programming experience, familiarity with JavaScript, and their opinion about the 10 atom candidates we explored in the both experiments. We first presented them pairs of clean and obfuscated code snippets

26

and then asked them to discuss their preferences towards one version. We did not indicate in any way whether any version was assumed to be confusing or not.

Table 11: Summary of participants' preferences for code snippets *with* and *without* atom candidates. The participants were only presented the code snippets, without any indication about whether one was confusing or not. Participants were also allowed to choose Neutral when they thought both sides were equally readable.

Atom	PREFERENCE (%)		
	Obfuscated	Clean	Neutral
Comma Operator	0	100	0
Automatic Semicolon Insertion	0	80	0
Post-Increment	20	73.33	6.67
Omitted Curly Braces	0	100	0
Assignment as Value	20	60	20
Implicit Predicate	20	73.33	6.67
Logic as Control Flow	20	60	20
Conditional Operator	60	26.67	13.3
Pre-Increment	40	46.67	13.33
Arithmetic as Logic	0	93.33	6.67
OVERALL	18	71.33	8.64

The Participants' perceptions of the atom candidates

Supporting the results of Sections 4 and 5, Table **??** shows that for eight out of the 10 scenarios surveyed, the majority of the respondents prefer the version of the code without the atom candidate. In no case the *neutral* ratio was higher than the option for the clean version. Only for the Conditional Operator atom candidate the participants preferred the obfuscated version, instead of the clean one. Figure 4 shows the code snippets for the corresponding obfuscated and clean versions. For this atom candidate, some participants who preferred the left-hand side version (obfuscated) still believed that the right-hand side version (clean) was more readable. The following quotes were extracted from the transcripts with two interviewees:

"I prefer to write [code using the left-hand side version], but I think [the right-hand side version] is easier to read, especially for newer programmers".

```
let config = {size: 3,
    isActive: false};
const_config = config.
    isActive === true
                ? config
                : {size: 10};
console.log(_config.size);
```

Listing 2: *Left-hand side* (using the *Conditional Operator* atom).

```
let config = {size: 3,
    isActive: false}
let _config;
if(config.isActive === true
    ) {
  _config = config;
}
else{
    _config = {size: 10};
}
console.log(_config.size);
```

Listing 3: *Right-hand side* (without the atom).

Figure 4: Example of a pair of code snippets used in the interview pertaining to the *Conditional Operator* atom candidate

> *"When I am programming, I write code with the conditional operator, [...], but, to be honest, I still think that [the code using the right-hand side version] is easier to understand"*.

The Pre-Increment atom candidate also caused a conflict in one of the interviewees, who regarded the clean version as simpler to understand, but would still opt to write code with the atom. In contrast, one of the participants found the version with the atom candidate more elegant, but recognized it was less readable, and was willing to sacrifice elegance for readability. In the repeated measures study, we found a significant difference in favor of the clean version, with a large effect size in its favor. In the latin square study, no difference could be observed in terms of correctness. However, there was a significant difference in terms of the time required to predict the output correctly in favor of the obfuscated version. This highlights how factors such as the experience of the participants and the selection of code snippets can have considerable impact on the results of studies about code readability [27].

(castor note) Should the last sentence above be in the discussion?

(rb note) Agree. We can move it to the discussion section.

When analyzing the Logic as Control Flow atom candidate, one of the interviewees gave an example of personal experience that might motivate one

28

to avoid writing code using this construct:

> *"This one is interesting, because I have written code that looks like the left-hand side [obfuscated version], and my colleagues complained that it was difficult to understand. Nowadays I prefer to write code using the version on the right-hand side [clean version]".*

None of the two studies found a significant difference in correctness between clean and obfuscated versions of snippets related to this atom candidate. However, in the latin square study the participants analyzing clean versions of the snippets required significantly less time.

For two of the atom candidates, the interviewees were unanimous in their preference for the clean version: Comma Operator and Omitted Curly Braces. Not coincidentally, in both experiments there were significant differences in correctness in favor of the clean versions for both atoms. Regarding the first one, we could often notice during the interviews that the *left-hand side* (with the atom of confusion) caused significant confusion among the participants. One remark about the *Comma Operator* atom of confusion is listed below:

> *"The code in [the left-hand side version] is unlikely to be understood unless the programmer knows C or C++".*

As for the atom candidate Omitted Curly Braces, one of the interviewees mentioned that, although they understand why one would opt not to use braces for simple if-then-else statements, they still advised against it, on grounds that:

> *"[I prefer the right-hand side version of the code ...] If I want to see well-written, easily understandable code, then I also have to do my job. Therefore I believe that, since I do not know who is on the other end maintaining this code, and it could be any person with any level of expertise, then I try to write readable, easy-to-understand code".*

Confusion in JavaScript Code

The final remarks drawn from the interviews are related to potentially confusing constructs suggested by the participants and their perspectives on JavaScript as a language. From the latter, we can also uncover some other forms in which the language itself might contribute to writing confusing code.

```
let inc = (x) => x + 1;
```

Listing 4: Example of arrow function.

```
let inc = function(x) {
   return x + 1
}
```

Listing 5: Alternative version without arrow function.

One of the participants mentioned that the use of JavaScript's prototype-based inheritance can make it difficult to understand code, particularly when involving deep prototype chains. When asked about particular JavaScript constructs or patterns that can make code difficult to understand, three participants cited the callback pattern—potentially leading to several levels of nested function calls—as extremely difficult to assimilate. One of the respondents stated:

> "Nested callbacks are very confusing. Even writing them can be confusing, let alone understanding them."

Two other developers implicitly touched upon the callback pattern, mentioning that it can be difficult to understand *asynchronous* programming in JavaScript. We also found other idioms that are JavaScript atom candidates, including *property access* via array subscription (`V1['P1']` instead of `V1.P1`) and *arrow functions* (see listings 4 and 5). The former was investigated in the repeated measures study but we did not find a significant difference between clean and obfuscated versions. As future work, we aim at investigating whether or not these atom candidates are more likely to introduce confusion in JavaScript code.

7. Results of the MSR Effort

In this section we present the results of the software repository mining effort, whose goal is to reveal how prevalent atom candidates are in open source JavaScript projects. We mined 72 open source JavaScript repositories to understand how often atoms arise in real software. As in the interview study, this study focused on atoms that were investigated in both experiments.

Similarly to previous studies [9, 6] which investigate the prevalence of atoms of confusion in open source C and C++ projects, we found that atoms

of confusion frequently arise in JavaScript open source systems. For instance, the six most frequently found atoms occur in at least 80% of the projects. Considering the extremes, atom candidates Conditional Operator and Implicit Predicate were found in all repositories we mined, while Comma Operator occurred in only 15% of them (as seen in Table ??).

Table 12: Summary of atom candidate occurrences in our dataset.

Atom	Projects	Occurr./KLOC
Implicit Predicate	100%	19.89
Conditional Operator	100%	10.16
Omitted Curly Braces	91.67%	6.61
Post Increment	90.28%	5.62
Pre Increment	83.33%	1.04
Assignment as Value	81.94%	1.02
Logic as Control Flow	56.94%	0.95
Automatic Semicolon Insertion	33.33%	0.13
Arithmetic as Logic	23.61%	0.02
Comma Operator	15.28%	0.03

Considering all JavaScript projects in our dataset, we found a total of 364873 atom candidates, though four atoms are responsible for 92.97% of this total: Implicit Predicate, Post Increment, Conditional Operator, and Omitted Curly Braces. The first two, based on the results of Section 5 and Section 4, can be considered atoms of confusion. The remaining atom candidates comprise 25621 occurrences combined, 7.03% of the total number of occurrences. The Comma Operator and Arithmetic as Logic atoms are the ones that arise less frequently (0.001% in total with 232 and 165 occurrences, respectively).

We calculate the frequency of each atom per KLOC. The four atoms mentioned in the previous paragraph, Implicit Predicate, Post Increment, Conditional Operator, and Omitted Curly Braces, occur frequently in the analyzed projects. All of them have more than 5 occurrences per KLOC on average. In particular, two atom candidates occur very frequently, Implicit Predicate (19.89 occurrences per KLOC) and Conditional Operator (10.16 occurrences per KLOC). On the other hand, Arithmetic as Logic and Comma Operator occur less than once for every 50 KLOC.

The results of our experiments and interviews suggest that Conditional Operator does not contribute significantly as a source of misunderstanding (increasing the number of wrong answers in 9% of the cases, according to the Latin-square study). Nonetheless, the Post-Increment Expression and

Automatic Semicolon Insertion atoms are listed in the top four sources of misunderstanding (Table 8), and also frequently appear in open source systems. As such, refactoring existing systems to avoid Post-Increment and Automatic Semicolon Insertion might improve the readability of large amounts of code in JavaScript projects.

8. Discussion

Comparing the results of the latin square study to previous research, three atoms of confusion that Gopstein et al. [4] confirmed for C and C++ do not lead to a statistically significant impact on correctness in our study: Logic as Control Flow, Conditional Operator, and Pre-Increment. Our results support the Gopstein et al. findings for the remaining atoms. The work of Langhout and Aniche [7] investigates six atom candidates for Java that we also explore here for JavaScript. Among these atom candidates, we achieve similar conclusions to their study for four atom candidates: Post Increment, Conditional Operator, Arithmetic as Logic, and Omitted Curly Braces—both studies confirmed only the former as an atom of confusion. Our findings diverged from the work of Langhout and Aniche for two atoms: Pre Increment and Logic as Control Flow.

> **(castor note)** Which atoms are JavaScript specific? Semicolons. Destructuring of objects and arrays.

> **(castor note)** Which atom candidates are atoms, considering these studies and the previous ones? Is a case where correctness exhibits no statistically significant difference but time does sufficient? The weight of experience. Statistical significance vs. practical relevance. Most of them, with the exception of Comma Operator and Object Destructuring, have also been observed for Java programs [7].
> In addition, many of the coding idioms that have been previously identified as atoms of confusion in previous work [4, 7], such as Logic as Control Flow
> Differences among the two studies, for example, post increment was not significant in the first one but it was significant in the second.

Our work has several implications. First, it adds external validity to the work of Gopstein et al. [4], which investigates the impact of atom candidates on understanding C,C++ code. That is, similarly to their work, the atom candidates Comma Operator, Post/Pre Increment, Omitted Curly Braces,

Assignment as Value, Implicit Predicate, and Logic as Control Flow seem to make JavaScript code harder to understand. For five of them, the difference is statistically significant, with a large effect size for four atoms. Our results also refute the hypothesis that Arithmetic as Logic is an atom of confusion (i.e., a source of misunderstanding). In comparison to the original work of Gopstein et al. [4], our study led to some differences in the effect size of the atom candidates. Altogether, we answer our first research question.

> **Answer to RQ1:** Our experiments using repeated measures and the Latin square designs give evidence that some of the atom candidates in C and C++ programs that also exist in JavaScript correspond to a source of misunderstanding in JavaScript code.

The results of the interview study complement the understanding of atoms of confusion because the participants make clear the existence of a trade-off between code comprehension and other quality attributes. For instance, most of the participants prefer the version of the code with the Conditional Operator, even though they agree that its use might contribute to the misunderstanding of JavaScript code, particularly when novices are maintaining the codebase. The participants of the interview study also mentioned other possible sources of misunderstanding in JavaScript, including the use of prototype-based inheritance and nested call-backs (as discussed in Section 6). Other JavaScript atom candidates include Object Destructuring, Array Spread, Object Spread, and Type Conversion. In summary, the results of the second study (interviews) allow us to answer the second research question.

> **Answer to RQ2:** The qualitative analysis of the interviews supports the results of our experiments— JavaScript developers most often agree that atoms of confusion compromise source code understanding.

The results of the third study (mining open source JavaScript repositories) provides evidence that, although atoms of confusion compromise program comprehension, they frequently appear in open source JavaScript projects. In particular, seven, out of 10 atoms considered in our study, appear in more than 50% of the 72 projects we analyzed. Furthermore, at least two of them are used intensively, more than once for every 200 lines of code. In summary, the third study allows us to answer the third research question.

Answer to RQ3: The MSR study reveals that several atom candidates explored in our research appear frequently in practice, and cleaning up the use of Post-Increment/Decrement and the Automatic Semicolon Insertion might improve the readabiliy of JavaScript code substantially.

9. Threats to Validity

In this section we discuss some of the main threats to the validity of this work.

Conclusion validity. In the context of our survey study, we apply different non-parametric statistical tests appropriate for the cases where data was categorical (correctness, Chi-square test of independence) and continuous (time, Mann-Whitney U test). Furthermore, besides reporting p-values, we also report effect size measures appropriate for each scenario (odds-ratio and Cliff's delta) and apply a p-value correction technique to avoid the multiple testing problem. Finally, it could be argued that the size of the samples is insufficient to make conclusions for some of the atoms, a common problem in Empirical Software Engineering. We estimate the sample size for each one of the atom candidates, considering that each one had a different number of samples (Table 8). To that end, we use the ϕ measure of effect size for each atom (based on the Chi-square statistic), the standard α coefficient of 0.05, and set the expected statistical power to 0.8, as usually employed in the literature [23]. We find out that the sample size is sufficient for all but one of the atoms where we had a statistically significant result for correctness, Implicit Predicate.

Construct validity. We use correctness as a proxy for program comprehension and predicted outcomes of small programs as a measure of correctness. As discussed elsewhere [28], this approach is a test of the developers' ability to trace programs. Although this is a common approach in studies about atoms of confusion [8, 7, 4], other measures of correctness could have been employed, potentially yielding different results. As a complement to correctness in the survey, we have measured the time it took for the participants to correctly predict the outcome of the code snippets (either with or without atom candidates). Furthermore, we have asked the interviewees about their preferences when comparing confusing and clean versions of the programs.

Internal validity. Since the survey was conducted online with unknown

34

participants, we have no way of confirming their levels of education and experience. We mitigate this threat by, in the data analysis, explicitly accounting for programmer experience and using an experimental design that allows us to isolate the impact of the treatment from factors such as experience and formal education level. Also, we did not have a way to prevent respondents from cheating, such as running the code on an interpreter, or consulting other people. We partially mitigate this threat by presenting images with source code, instead of text. This creates an obstacle for participants to run the code while taking the survey.

External validity. Our results suggest that the selected idioms and code constructs may lead to confusion for small code snippets, but it is not clear if that result extrapolates to other scenarios. Even though it is likely that in larger code bases the confusion induced by these constructs may be even greater, the existence of additional context may mitigate this effect. Another possible threat to the generalizability of our results, in particular for the survey and interviews, lies in the fact that the analyzed atoms may rarely occur in real JavaScript code bases. To mitigate that threat, we have analyzed 72 popular open source JavaScript repositories and found out that most of them are common, occurring at least once per 1,000 lines of code. Only Arithmetic as Logic and Comma Operator occur less frequently than once per 10,000 lines of code.

10. Conclusions

This paper reports the results of a mixed-method research effort that allowed us to better understand the impact of atoms of confusion in JavaScript code. First, we conducted a survey with 140 JavaScript developers, asking them to predict the output of some code snippets (with and without atoms). We provide evidence that five atom candidates significantly impact program comprehension activities and can be considered atoms of confusion. Our efforts have two implications for practice. First, we present evidence that not all atoms of confusion validated to statically typed languages (such as C, C++, or Java) led to a statistically significant impact on program understanding for JavaScript (a dynamically typed language). Besides that, our findings might be used to alert developers to avoid writing JavaScript code with certain atoms of confusion (e.g., Comma Operator, Automatic Semicolon Insertion, Post Increment/Decrement, Assignment as Value, and Implicit Predicate). Finally, our results might help tool developers to cre-

ate program transformation tools for removing atoms that frequently appear in software. Our research also pointed out additional JavaScript constructs and idioms that might introduce misunderstanding. These include nested callbacks, *property access*, and *arrow functions*. As future work, we intend to reproduce our survey to validate whether or not these additional *atom candidates* truly affect the understanding of JavaScript code.

References

[1] F. Ebert, F. Castor, N. Novielli, A. Serebrenik, An exploratory study on confusion in code reviews, Empir. Softw. Eng. 26 (2021) 12.

[2] F. Hermans, The Programmer's Brain: What every programmer needs to know about cognition, Manning, 2021.

[3] S. Ajami, Y. Woodbridge, D. G. Feitelson, Syntax, predicates, idioms: what really affects code complexity?, in: Proceedings of the 25th International Conference on Program Comprehension, ICPC, IEEE Computer Society, Buenos Aires, Argentina, 2017, pp. 66–76.

[4] D. Gopstein, J. Iannacone, Y. Yan, L. DeLong, Y. Zhuang, M. K. Yeh, J. Cappos, Understanding misunderstandings in source code, in: ESEC/SIGSOFT FSE, ACM, https://doi.org/10.1145/3106237.3106264, 2017, pp. 129–139.

[5] F. Castor, Identifying confusing code in swift programs, In Proceedings of the VI CBSoft Workshop on Visualization, Evolution, and Maintenance 1 (2018) 1–8.

[6] F. Medeiros, G. Lima, G. Amaral, S. Apel, C. Kästner, M. Ribeiro, R. Gheyi, An investigation of misunderstanding code patterns in C open-source software projects, Empirical Software Engineering 24 (2019) 1693–1726.

[7] C. Langhout, M. Aniche, Atoms of confusion in java, in: Proceedings of the 29th IEEE/ACM International Conference on Program Comprehension, ICPC '21, New York, NY, USA, 2021. To appear.

[8] B. de Oliveira, M. Ribeiro, J. A. S. da Costa, R. Gheyi, G. Amaral, R. de Mello, A. Oliveira, A. Garcia, R. Bonifácio, B. Fonseca, Atoms of confusion: The eyes do not lie, in: Proceedings of the 34th Brazilian Symposium on Software Engineering, SBES '20, Association for Computing Machinery, New York, NY, USA, 2020, p. 243–252. URL: `https://doi.org/10.1145/3422392.3422437`. doi:10.1145/3422392.3422437.

[9] D. Gopstein, H. H. Zhou, P. G. Frankl, J. Cappos, Prevalence of confusing code in software projects: atoms of confusion in the wild, in: MSR, ACM, https://doi.org/10.1145/3196398.3196432, 2018, pp. 281–291.

[10] B. Eich, A brief history of JavaScript, 2018. Available at https://www.youtube.com/watch?v=GxouWy-ZE80. Last access: July 25th 2022.

[11] H. J. Keselman, J. Algina, R. K. Kowalchuk, The analysis of repeated measures designs: A review, British Journal of Mathematical and Statistical Psychology 54 (2001) 1–20.

[12] E. George, W. G. Hunter, J. S. Hunter, Statistics for experimenters: Design, innovation, and discovery, Wiley, 2005.

[13] S. R. Tilley, D. B. Smith, S. Paul, Towards a framework for program understanding, in: WPC, IEEE Computer Society, DOI: 10.1109/WPC.1996.501117, 1996, p. 19.

[14] A. B. O'Hare, E. W. Troan, Re-analyzer: From source code to structured analysis, IBM Systems Journal 33 (1994) 110–130.

[15] D. Gopstein, A.-L. Fayard, S. Apel, J. Cappos, Thinking aloud about confusing code: A qualitative investigation of program comprehension and atoms of confusion, ESEC/FSE 2020, Association for Computing Machinery, New York, NY, USA, 2020, p. 605–616. URL: https://doi.org/10.1145/3368089.3409714. doi:10.1145/3368089.3409714.

[16] F. Alves, D. Oliveira, F. Madeiral, F. Castor, On the bug-proneness of structures inspired by functional programming in javascript projects, CoRR abs/2206.08849 (2022). URL: https://doi.org/10.48550/arXiv.2206.08849.

[17] J. H. Neely, Semantic priming effects in visual word recognition: A selective review of current findings and theories, Lawrence Erlbaum Associates, Inc., 1991, p. 264–336.

[18] P. Gyimesi, B. Vancsics, A. Stocco, D. Mazinanian, A. Beszédes, R. Ferenc, A. Mesbah, Bugsjs: A benchmark of javascript bugs, in: 2019 12th IEEE Conference on Software Testing, Validation and Verification (ICST), IEEE, 2019, pp. 90–101.

[19] E. D. Canedo, R. Bonifácio, M. V. Okimoto, A. Serebrenik, G. Pinto, E. Monteiro, Work practices and perceptions from women core developers in OSS communities, in: ESEM '20: ACM / IEEE International Symposium on Empirical Software Engineering and Measurement, Bari, Italy, October 5-7, 2020, ACM, 2020, pp. 26:1–26:11. URL: `https://doi.org/10.1145/338249 4.3410682`. doi:10.1145/3382494.3410682.

[20] O. de Moor, D. Sereni, M. Verbaere, E. Hajiyev, P. Avgustinov, T. Ekman, N. Ongkingco, J. Tibble, .ql: Object-oriented queries made easy, in: R. Lämmel, J. Visser, J. Saraiva (Eds.), Generative and Transformational Techniques in Software Engineering II, International Summer School, GTTSE 2007, Braga, Portugal, July 2-7, 2007. Revised Papers, volume 5235 of *Lecture Notes in Computer Science*, Springer, 2007, pp. 78–133. URL: `https://doi.org/10.1007/978-3-540-88643-3\_3`. doi:10.1007/978-3-540-88643-3_3.

[21] O. Rodriguez-Prieto, F. Ortin, An efficient and scalable platform for java source code analysis using overlaid graph representations (support material website), 2020.

[22] Q. McNemar, Note on the sampling error of the difference between correlated proportions or percentages, Psychometrika 12 (1947) 153—-157.

[23] P. D. Ellis, The Essential Guide to Effect Sizes: Statistical Power, Meta-Analysis, and the Interpretation of Research Results, Cambridge University Press, 2010.

[24] S. S. Mangiafico, Summary and Analysis of Extension Program Evaluation in R, New Brunswick: Rutgers Cooperative Extension, 2016. Available at rcompanion.org/handbook/.

[25] H. Chen, P. Cohen, S. Chen, How big is a big odds ratio? interpreting the magnitudes of odds ratios in epidemiological studies, Communications in Statistics - Simulation and Computation 39 (2010) 860–864.

[26] J. Romano, J. Kromrey, J. Coraggio, J. Skowronek, Appropriate statistics for ordinal level data: Should we really be using t-test and Cohen'sd for evaluating group differences on the NSSE and other surveys?, in: Proceedings of the Annual meeting of the Florida Association of Institutional Research, 2006, pp. 1–3.

[27] D. G. Feitelson, Considerations and pitfalls in controlled experiments on code comprehension, in: 29th IEEE/ACM International Conference on Program Comprehension, IEEE, Madrid, Spain, 2021, pp. 106–117.

[28] D. Oliveira, R. Bruno, F. Madeiral, F. Castor, Evaluating code readability and legibility: An examination of human-centric studies, in: Proceedings of the IEEE International Conference on Software Maintenance and Evolution, IEEE, Adelaide, Australia, 2020, pp. 348–359.